

Homework 4 – Generative AI (CS/DS 552, Whitehill, Spring 2025)

Collaboration policy: You may complete this assignment with a partner, if you choose. In that case, both partners should sign up on Canvas in a **pre-made group** as a team, and only one of you should submit the assignment. You are permitted to use ChatGPT (or another AI tool) without restriction to help save you time typing boilerplate code, making complicated visualizations, and overcoming tedious syntactic issues. However, *you must fully understand all the work that you submit.*

1. **Diffusions: Unconditional Generation [13 pts]:** Using a trained diffusion model (you can either just use the pre-trained DDPM in `diff_unet_faces.cpt` or, if you prefer, train your own using the face data in `faces23k_48x48.npy.gz`), sample 100 faces (each is 48x48 pixels) and display them in a 10x10 collage. For DDPM, the sampling procedure is given in the Prince textbook (linked on Canvas), Algorithm 18.2. For σ_t , a common practice (according to the original DDPM paper by Ho et al, 2020) is to set $\sigma_t = \sqrt{\beta_t}$. The hyperparameters used for training, including the noise schedule $\{\beta_t\}_{t=1}^T$, are given in the starter code `homework4_template.py`.
2. **Diffusions: Inpainting/Merging [12 pts]:** Using the same pre-trained DDPM from the previous problem, “merge” two faces into one:
 - (a) Select any two faces $\mathbf{x}^{(1)}$ and $\mathbf{x}^{(2)}$.
 - (b) Apply the diffusion kernel for some $t < T$ to compute $\mathbf{z}_t^{(1)}$ and $\mathbf{z}_t^{(2)}$.
 - (c) Create \mathbf{z}_t by combining the left half of $\mathbf{z}_t^{(1)}$ with the right half of $\mathbf{z}_t^{(2)}$.
 - (d) Sample $\mathbf{x} \sim P(\mathbf{x} | \mathbf{z}_t)$.

Apply this to 5 pairs of faces each for 4 different values of t each – show the results in a 4x15 collage (where the rows correspond to different t – 10 columns show the 5 pairs of $(\mathbf{x}^{(1)}, \mathbf{x}^{(2)})$, and 5 columns show the corresponding merged samples \mathbf{x} . Arrange the columns as you see fit but make sure to explain them clearly in your report.

3. **Diffusions: Conditional Generation with Classifier Guidance [30 pts]:** Using the face data in `faces23k_48x48.npy` and associated age labels in `ages23k.npy`, train an age regressor f that maps from any (\mathbf{z}_t, t) to an estimate \hat{y} of how old the person in the corresponding \mathbf{x} is. To create the training data for f , you will need to use the diffusion kernel to create \mathbf{z}_t for various t and \mathbf{x} . Exclude any examples from training if either $y < 0$ (which indicates a missing label) or $y > 100$ (which is likely an incorrect label). f should output the mean of a Gaussian distribution with unit variance that characterizes the age of a specific noised image at a specific timestep t , i.e., $P(y | \mathbf{z}_t) = \mathcal{N}(f(\mathbf{z}_t, t), 1)$. As an architecture for the regressor, you can use the first half (i.e., downward convolution) of the same U-Net with time embeddings which implements the noise estimator E_θ in the pre-trained DDPM. However, the choice is up to you.

After training the regressor, use classifier guidance (even though you actually have a regressor) to generate 4 faces each for target ages 18, 40, 60, and 80 years old. To compute $\frac{\partial \log P(y | \mathbf{z}_t)}{\partial \mathbf{z}_t} = -\frac{\partial (\hat{y} - y)^2}{\partial \mathbf{z}_t}$ for a target age value y , you can ask ChatGPT for help with the autograd expressions (you will need to enable `requires_grad=True` on the tensor representing \mathbf{z}_t). You may find it useful to multiply the gradient by a small number (e.g., 0.2) rather than using its full magnitude. Show these faces in a 4x10 collage.

Submission: Create a Zip file containing all your Python code (which can be in multiple files) and your PDF file, and then submit on Canvas. If you are working as part of a group, then only **one** member of your group should submit (but make sure you have already signed up in a pre-allocated team for the homework on Canvas).